

1 Dijkstra's Algorithm & Priority Queues

Given a graph where each edge has some length ℓ , how do we find the shortest path between two vertices?

1.1 Trivial Way

The idea is to break edges into unit length edges. So, an edge of length 5 can actually be seen as 5 edges of length 1. With this conversion, we can run BFS. However, the runtime is $O(\text{Sum of Edge Lengths})$.

1.2 Another Way

If we have very long edge lengths, most steps will just consist of advancing slightly along a bunch of edges. So, how do we “fast forward” through these boring steps? Well, occasionally, we have interesting steps where the wavefront hits a new vertex.

1.3 Algorithm

```

Distances(G, s, l)
  dist(s) = 0
  While not all distances found
    Find minimum over (v, w) in E
      with v discovered w not
      of dist(v) + l(v, w)
    dist(w) = dist(v) + l(v, w)
    prev(w) = v

```

1.4 Why Does This Work?

Proposition. *Whenever the algorithm assigns a distance to a vertex v , that is the length of the shortest path from s to v .*

Proof. We use induction.

Base Case: We know that $\text{dist}(s) = 0$. The empty path has length 0.

Inductive Step: When assigning distances to w , suppose that all previously assigned distances are correct. TODO

□

This runs in $O(|V| \cdot |E|)$ time. There are $O(|V|)$ iterations and $O(|E|)$ edges.

- This is too slow because every iteration we have to check every edge.
- The idea is that most of the comparison doesn't change much iteration to iteration. So, we can use this to save time.
- Specifically, record for each w the best value of $\text{dist}(v, w) + \ell(v, w)$.

1.5 Better Algorithm

```

Distances(G, s, l)
  For v in V
    dist(v) = infinity
    done(v) = false
  dist(s) = 0
  done(s) = false

```

```

while not all vertices done
    Find v not done with minimum dist(v)
    done(v) = true
    For (v, w) in E
        if dist(v) + l(v, w) < dist(w)
            dist(w) = dist(v) + l(v, w)
            prev(w) = v

```

The initialization is $O(|V|)$, and the while loop is $O(|V|)$. The for loop is $O(|E|)$ time. Thus, the runtime is:

$$O(|V|^2 + |E|)$$

- This repeatedly asks for the smallest vertex. Even though not much is changing from round to round, the algorithm is computing the minimum from scratch every time.
- We can use a data structure to help answer a bunch of similar questions faster than answering each question individually (like the one above).

1.6 Priority Queue

A **priority queue** is a data structure that stores elements sorted by a **key** value. Its operations are:

- **insert**: Adds a new element to the priority queue.
- **decreaseKey**: Changes the key of an element of the priority queue to a specified smaller value.
- **deleteMin**: Deletes the element with the lowest key from the priority queue.

1.7 Even Better Priority Queue

```

Dijkstra(G, s, l)
    Initialize Priority Queue Q
    For v in V
        dist(v) = infinity
        Q.insert(v) // using dist(v) as key
    dist(s) = 0
    while Q not empty
        v = Q.deleteMin()
        For (v, w) in E
            if dist(v) + l(v, w) < dist(w)
                dist(w) = dist(v) + l(v, w)
                prev(w) = v
                // w has a faster path and needs to be updated in
                // the priority queue
                Q.decreaseKey(w)

```

The runtime is as follows:

- We need to iterate $O(|V|)$ times for the initial loop.
- In the **while** loop, we run $O(|V|)$ times.
- We need to run through the edges $O(|E|)$ times.

So, $O(|V|)$ inserts + $O(|V|)$ deleteMins + $O(|E|)$ decreaseKey.